

✓ Telecom Customer Churn Analysis

Customer Churn is one of the most important metrics for a business to track. Customer acquisition costs are a massive factor in scaling. Ultimately a successful business model will hinge upon preventing customer churn.



Data Analysis and Machine Learning are exceptional tools for understanding Churn. Data Analysis aids in understanding trends while machine learning enables the creation of predictive models. By combining these tools with strategic data collection a business model can be iteratively updated to ensure continuous revenue growth and scalability.

✓ Exploratory Data Analysis and Visualization

✓ Wrangling

Our first steps are to ingest the data, gain an understanding of its structure, and the distribution of its values

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
telecom_data = pd.read_csv('/content/WA_Fn-UseC_-Telco-Customer-Churn.csv')
```

```
telecom_data.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	N
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	N
4	9237-HQITU	Female	0	No	No	2	Yes	

5 rows × 21 columns

✓ Columns

```
telecom_data.columns
```

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

```
telecom_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
#   ...
```




```
--- -----
0 customerID 7043 non-null object
1 gender 7043 non-null object
2 SeniorCitizen 7043 non-null int64
3 Partner 7043 non-null object
4 Dependents 7043 non-null object
5 tenure 7043 non-null int64
6 PhoneService 7043 non-null object
7 MultipleLines 7043 non-null object
8 InternetService 7043 non-null object
9 OnlineSecurity 7043 non-null object
10 OnlineBackup 7043 non-null object
11 DeviceProtection 7043 non-null object
12 TechSupport 7043 non-null object
13 StreamingTV 7043 non-null object
14 StreamingMovies 7043 non-null object
15 Contract 7043 non-null object
16 PaperlessBilling 7043 non-null object
17 PaymentMethod 7043 non-null object
18 MonthlyCharges 7043 non-null float64
19 TotalCharges 7043 non-null object
20 Churn 7043 non-null object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

Values

```
unique_values = telecom_data.apply(lambda col: col.unique())

unique_values = pd.DataFrame(unique_values)

unique_values
```

0		
customerID	[7590-VHVEG, 5575-GNVDE, 3668-QPYBK, 7795-CFOC...	
gender	[Female, Male]	
SeniorCitizen	[0, 1]	
Partner	[Yes, No]	
Dependents	[No, Yes]	
tenure	[1, 34, 2, 45, 8, 22, 10, 28, 62, 13, 16, 58, ...]	
PhoneService	[No, Yes]	
MultipleLines	[No phone service, No, Yes]	
InternetService	[DSL, Fiber optic, No]	
OnlineSecurity	[No, Yes, No internet service]	
OnlineBackup	[Yes, No, No internet service]	
DeviceProtection	[No, Yes, No internet service]	
TechSupport	[No, Yes, No internet service]	
StreamingTV	[No, Yes, No internet service]	
StreamingMovies	[No, Yes, No internet service]	
Contract	[Month-to-month, One year, Two year]	
PaperlessBilling	[Yes, No]	
PaymentMethod	[Electronic check, Mailed check, Bank transfer...	
MonthlyCharges	[29.85, 56.95, 53.85, 42.3, 70.7, 99.65, 89.1, ...]	
TotalCharges	[29.85, 1889.5, 108.15, 1840.75, 151.65, 820.5...	
Churn	[No, Yes]	

```
telecom_data.isnull().sum()

customerID      0
gender          0
SeniorCitizen   0
Partner         0
```

```

Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64

```

Currently there are no null values in the data, but some of the data is of the wrong type which we will need correct in order to perform regression. I will also drop Customer ID because it is not a feature related to Churn. I suspect that the variables 'PhoneService' and 'MultipleLines' are confounding. During regression analysis I may have to address multicollinearity.

```
telecom_data['TotalCharges'] = pd.to_numeric(telecom_data['TotalCharges'], errors='coerce')
```

```
telecom_data['SeniorCitizen'] = telecom_data['SeniorCitizen'].astype(object).replace({0: 'No', 1: 'Yes'})
```

```
telecom_data.drop(columns = 'customerID', inplace=True)
```

```
telecom_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   gender                 7043 non-null  object
1   SeniorCitizen          7043 non-null  object
2   Partner                7043 non-null  object
3   Dependents             7043 non-null  object
4   tenure                 7043 non-null  int64
5   PhoneService           7043 non-null  object
6   MultipleLines          7043 non-null  object
7   InternetService        7043 non-null  object
8   OnlineSecurity         7043 non-null  object
9   OnlineBackup           7043 non-null  object
10  DeviceProtection       7043 non-null  object
11  TechSupport            7043 non-null  object
12  StreamingTV            7043 non-null  object
13  StreamingMovies        7043 non-null  object
14  Contract               7043 non-null  object
15  PaperlessBilling       7043 non-null  object
16  PaymentMethod          7043 non-null  object
17  MonthlyCharges         7043 non-null  float64
18  TotalCharges           7032 non-null  float64
19  Churn                  7043 non-null  object
dtypes: float64(2), int64(1), object(17)
memory usage: 1.1+ MB

```

```
telecom_data.isnull().sum()
```

```

gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0

```

```

PaymentMethod      0
MonthlyCharges     0
TotalCharges       11
Churn              0
dtype: int64

```

```
telecom_data.loc[telecom_data['TotalCharges'].isnull()]
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	:
488	Female	No	Yes	Yes	0	No	No phone service	
753	Male	No	No	Yes	0	Yes	No	
936	Female	No	Yes	Yes	0	Yes	No	
1082	Male	No	Yes	Yes	0	Yes	Yes	
1340	Female	No	Yes	Yes	0	No	No phone service	
3331	Male	No	Yes	Yes	0	Yes	No	
3826	Male	No	Yes	Yes	0	Yes	Yes	
4380	Female	No	Yes	Yes	0	Yes	No	
5218	Male	No	Yes	Yes	0	Yes	No	
6670	Female	No	Yes	Yes	0	Yes	Yes	
6754	Male	No	No	Yes	0	Yes	Yes	

The data wrangling process has created some Null values, after inspecting them we can see other issues with these entries as well. For instance the customer's have monthly charges but no tenure. Is this their first month with the service? I will inspect what other customer's have a tenure of 0 months.

```
telecom_data.loc[telecom_data['tenure']==0]
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	:
488	Female	No	Yes	Yes	0	No	No phone service	
753	Male	No	No	Yes	0	Yes	No	
936	Female	No	Yes	Yes	0	Yes	No	
1082	Male	No	Yes	Yes	0	Yes	Yes	
1340	Female	No	Yes	Yes	0	No	No phone service	
3331	Male	No	Yes	Yes	0	Yes	No	
3826	Male	No	Yes	Yes	0	Yes	Yes	
4380	Female	No	Yes	Yes	0	Yes	No	
5218	Male	No	Yes	Yes	0	Yes	No	
6670	Female	No	Yes	Yes	0	Yes	Yes	
6754	Male	No	No	Yes	0	Yes	Yes	

It appears they are all the same customer. I could impute the values using the mean of 'TotalCharges' or a classification algorithm, though in this situation it makes the most sense to just replace 'TotalCharges' with 0 as I believe they are first month customers.

```
telecom_data['TotalCharges'].fillna(0, inplace=True)
```

Now that the data has the correct variable types and there are no null values, I will separate the label and the features by type. I also want to standardize the naming of the variables.

```
telecom_data.columns = telecom_data.columns.str.capitalize()
telecom_data.columns
```

```
Index(['Gender', 'Seniorcitizen', 'Partner', 'Dependents', 'Tenure',
       'Phoneservice', 'Multiplelines', 'Internetservice', 'Onlinesecurity',
       'Onlinebackup', 'Deviceprotection', 'Techsupport', 'Streamingtv',
       'Streamingmovies', 'Contract', 'Paperlessbilling', 'Paymentmethod',
       'Monthlycharges', 'Totalcharges', 'Churn'],
      dtype='object')
```


```
label = 'Churn'
categorical_features = []
numerical_features = []
```

```
for column in telecom_data.columns:
    if column == label:
        pass
    elif telecom_data[column].dtype == 'object':
        categorical_features.append(column)
    else:
        numerical_features.append(column)
```

```
print("Categorical features:", categorical_features)
print("Numerical features:", numerical_features)
print("Label:", label)
```

```
Categorical features: ['Gender', 'Seniorcitizen', 'Partner', 'Dependents', 'Phoneservice', 'Multiplelines', 'Internetservice', 'Onlines
Numerical features: ['Tenure', 'Monthlycharges', 'Totalcharges']
Label: Churn
```

```
telecom_data[label].value_counts().plot(kind='bar')
plt.show()
```



```
churn_percentage = (telecom_data['Churn'].value_counts(normalize=True)* 100)['Yes']
print(f"Churn Percentage: {churn_percentage:.2f}%")
```

```
Churn Percentage: 26.54%
```

Over 25% of customers have churned, this is much higher than we would like. I will now examine and visualize the features of this dataset to see if there are features that are causing this churn rate to be so high.

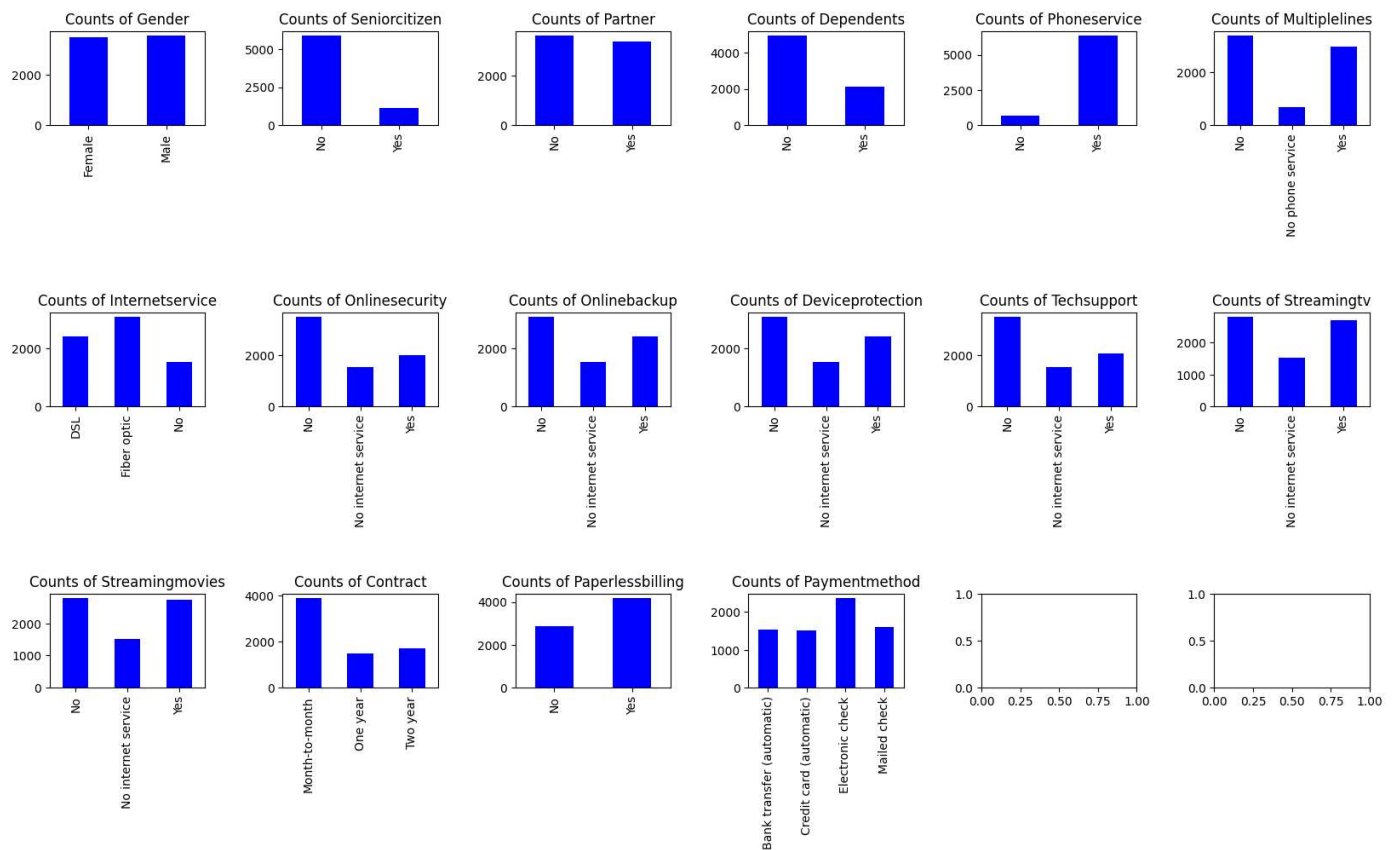
Visualization

Categorical Data

```
fig, axes = plt.subplots(3,6, figsize=(20, 10), gridspec_kw={'wspace': 0.5, 'hspace':2})
axes = axes.flatten()

# Plot counts for each categorical feature
for i, column in enumerate(categorical_features):
    telecom_data[column].value_counts().sort_index().plot(kind='bar', color='blue', ax=axes[i])
    axes[i].set_title(f'Counts of {column}')

plt.show()
```




Key Takeaways

- Gender is equally distributed
- The majority of customers are not seniors
- Most members have phone and internet service
- Most customers have a month-month contract

Numerical Data


```
telecom_data[numerical_features].hist(color = 'pink', edgecolor='blue', bins=30,
                                     figsize = (20,10), layout=(2,4))
plt.show()
```



```
fig, axes = plt.subplots(1, 3, figsize=(15, 5), gridspec_kw={'wspace': 0.5})

# Plot each numerical feature in a separate subplot
for i, feature in enumerate(numerical_features):
    telecom_data[feature].plot(kind='box', vert=False, ax=axes[i], color= 'blue')
    axes[i].set_title(f'Box Plot for {feature}')

plt.show()
```



```
telecom_data[numerical_features].describe()
```

```

    Tenure  Monthlycharges  Totalcharges
telecom_data.groupby(['Churn']).agg({'Monthlycharges': 'mean', 'Totalcharges': 'mean', 'Tenure': 'mean'})

```

	Monthlycharges	Totalcharges	Tenure
Churn			
No	61.265124	2549.911442	37.569965
Yes	74.441332	1531.796094	17.979133
75%	55.000000	89.850000	37.86.600000

Key Takeaways

- The average customer pays \$65 a month for service
 - They have been customers for 30 months
- The second highest category of all customers is those who have over 60 months of tenure
- Customers that churn pay slightly more a month on average
 - Because they churn this reduces their total charges
- The average churning customer will have service for 17 months

Relationship Between Numeric Variables

```
telecom_data.corr(numeric_only=True)
```

	Tenure	Monthlycharges	Totalcharges
Tenure	1.000000	0.247900	0.826178
Monthlycharges	0.247900	1.000000	0.651174
Totalcharges	0.826178	0.651174	1.000000

This is straightforward, Total Charges is highly correlated with Monthly Charges and Tenure

Relationship Between Categorical Variables

```

table_cnt = telecom_data.groupby(['Churn', 'Internetservice']).\
agg(cnt = ('Gender', lambda x: len(x)))
print(table_cnt)

```

Churn	Internetservice	cnt
No	DSL	1962
	Fiber optic	1799
	No	1413
Yes	DSL	459
	Fiber optic	1297
	No	113

```

table_cnt = telecom_data.groupby(['Churn', 'Seniorcitizen']).\
agg(cnt = ('Gender', lambda x: len(x)))
print(table_cnt)

```

Churn	Seniorcitizen	cnt
No	No	4508
	Yes	666
Yes	No	1393
	Yes	476

Because most of our data is categorical we will move onto **training a model**.

This will move the problem forward in two ways:

1. Creating a model which can be used on future data
2. Training a model to identify which features are most important in predicting churn

Machine Learning

Model Training and Evaluation

```

from sklearn.preprocessing import OneHotEncoder, LabelEncoder, StandardScaler
from sklearn.compose import make_column_transformer

# categories

category_transformer = make_column_transformer(
    #(OneHotEncoder(drop='first'), categorical_features))
    (OneHotEncoder(), categorical_features))

category_transformed = category_transformer.fit_transform(telecom_data[categorical_features])

ohe_df = pd.DataFrame(category_transformed, columns=category_transformer.get_feature_names_out(categorical_features))

# numerical

numerical_transformer = make_column_transformer(
    (StandardScaler(), numerical_features))

numerical_transformed = numerical_transformer.fit_transform(telecom_data[numerical_features])

scaled_df = pd.DataFrame(numerical_transformed, columns=numerical_transformer.get_feature_names_out(numerical_features))

# join dataframes

transformed_df = ohe_df.join(scaled_df)

# label

le = LabelEncoder()
label_df = pd.DataFrame(le.fit_transform(telecom_data[label]), columns = [label])

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, ConfusionMatrixDisplay

X_train, X_test, y_train, y_test = train_test_split(transformed_df, label_df, test_size=0.25, random_state=42)

# Logistic Regression Model
logistic_model = LogisticRegression()
logistic_model.fit(X_train, y_train)
y_pred_logistic = logistic_model.predict(X_test)
accuracy_logistic = accuracy_score(y_test, y_pred_logistic)
print("Accuracy of Logistic Regression:", accuracy_logistic)

# Random Forest Model
random_forest_model = RandomForestClassifier(n_estimators=100)
random_forest_model.fit(X_train, y_train)
y_pred_random_forest = random_forest_model.predict(X_test)
accuracy_random_forest = accuracy_score(y_test, y_pred_random_forest)
print("Accuracy of Random Forest:", accuracy_random_forest)

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result(
<ipython-input-31-7b492f6ea878>:17: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
random_forest_model.fit(X_train, y_train)
Accuracy of Logistic Regression: 0.8126064735945485
Accuracy of Random Forest: 0.787052810902896

```


```
cm = confusion_matrix(y_test, y_pred_logistic)

ConfusionMatrixDisplay(confusion_matrix=cm).plot()

lgc_accuracy = accuracy_score(y_test, y_pred_logistic)
lgc_precision = precision_score(y_test, y_pred_logistic)
lgc_recall = recall_score(y_test, y_pred_logistic)

print("Accuracy:", lgc_accuracy)
print("Precision:", lgc_precision)
print("Recall:", lgc_recall)
```

```
Accuracy: 0.8126064735945485
Precision: 0.6839506172839506
Recall: 0.5782881002087683
```




```
cm = confusion_matrix(y_test, y_pred_random_forest)

ConfusionMatrixDisplay(confusion_matrix=cm).plot()

rndf_accuracy = accuracy_score(y_test, y_pred_random_forest)
rndf_precision = precision_score(y_test, y_pred_random_forest)
rndf_recall = recall_score(y_test, y_pred_random_forest)

print("Accuracy:", rndf_accuracy)
print("Precision:", rndf_precision)
print("Recall:", rndf_recall)
```


```
Accuracy: 0.787052810902896
Precision: 0.6595092024539877
Recall: 0.4488517745302714
```



```
# Feature importance
feature_importances = pd.Series(random_forest_model.feature_importances_, index=X_train.columns).sort_values(ascending=False)

# Plot a simple bar chart
feature_importances = feature_importances[:10]

feature_importances.plot.bar()
```




▼ Oversampling Technique

Because Churn is imbalanced, the accuracy of the model can be improved by synthetically oversampling the minority value.

```
sns.countplot(x = telecom_data[label])
```


<Axes: xlabel='Churn', ylabel='count'>



```
from imblearn.over_sampling import SMOTE
X_res,y_res = SMOTE().fit_resample(transformed_df,label_df)
```

```
sns.countplot(x = y_res['Churn'])
```

<Axes: xlabel='Churn', ylabel='count'>



```
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.25, random_state=42)
```

```
# Logistic Regression Model
```

```
smote_logistic_model = LogisticRegression()
smote_logistic_model.fit(X_train, y_train)
y_pred_smote_logistic = smote_logistic_model.predict(X_test)
accuracy_smote_logistic = accuracy_score(y_test, y_pred_smote_logistic)
print("Accuracy of Logistic Regression:", accuracy_smote_logistic)
```

```
# Random Forest Model
```

```
smote_random_forest_model = RandomForestClassifier(n_estimators=100)
smote_random_forest_model.fit(X_train, y_train)
y_pred_smote_random_forest = smote_random_forest_model.predict(X_test)
accuracy_smote_random_forest = accuracy_score(y_test, y_pred_smote_random_forest)
print("Accuracy of Random Forest:", accuracy_smote_random_forest)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d
y = column_or_1d(y, warn=True)
Accuracy of Logistic Regression: 0.7943563973714728
<ipython-input-38-6e0fe16d0ab6>:12: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
smote_random_forest_model.fit(X_train, y_train)
Accuracy of Random Forest: 0.8565906455353691
```


```
cm = confusion_matrix(y_test, y_pred_smote_random_forest)
```

```
ConfusionMatrixDisplay(confusion_matrix=cm).plot()
```

```
smote_rndf_accuracy = accuracy_score(y_test, y_pred_smote_random_forest)
smote_rndf_precision = precision_score(y_test, y_pred_smote_random_forest)
smote_rndf_recall = recall_score(y_test, y_pred_smote_random_forest)
```

```
print("Accuracy:", smote_rndf_accuracy)
print("Precision:", smote_rndf_precision)
print("Recall:", smote_rndf_recall)
```

Accuracy: 0.8565906455353691
 Precision: 0.86198738170347
 Recall: 0.847944142746315




```
# Feature importance
feature_importances = pd.Series(smote_random_forest_model.feature_importances_, index=X_train.columns).sort_values(ascending=False)

# Plot a simple bar chart
feature_importances = feature_importances[:10]

feature_importances.plot.bar()
```

<Axes: >



✓ Conclusion

It should be no surprise that tenure is the most important feature for predicting churn, as churn directly affects tenure.

What is more important to recognize is that the month to month contract also predicts churn.

It is these customers that should be focused on as they are uncommitted and most likely to churn. The company should focus on keeping these customers past the average of 30 months.

The Random Forest Model which was trained using the minority oversampling technique can be used on future customers to predict their likelihood of churning.